

HarvardX: PH125.9x Data Science – Capstone Project: MovieLens

Project by: Eshan Sama

25th December 2020

Introduction

This capstone project creating a movie recommendation system using the MovieLens dataset. Also to predicts the movie rating by machine learning algorithm using the inputs in one subset in the validation set. The dataset used for this purpose can be found in the following links

- [MovieLens 10M dataset] <https://grouplens.org/datasets/movielens/10m/>
- [MovieLens 10M dataset - zip file] <http://files.grouplens.org/datasets/movielens/ml-10m.zip>

The goal of this project is:

- To create a movie recommendation system
- To reduce the RMSE less than 0.8649

Read the data

```
library(tidyverse)

library(caret)

dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)

ratings <- read.table(text = gsub("::", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
                      col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\: ", 3)
colnames(movies) <- c("movieId", "title", "genres")
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
                                           title = as.character(title),
                                           genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")
```

```
# Validation set will be 10% of MovieLens data
```

```
set.seed(1)
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, l
ist = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]
```

```
# Make sure userId and movieId in validation set are also in edx set
```

```
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")
```

```
# Add rows removed from validation set back into edx set
```

```
removed <- anti_join(temp, validation)
```

```
## Joining, by = c("userId", "movieId", "rating", "timestamp", "title", "genres")
```

DATA EXPLORATION

```
summary(edx)
```

```
##      userId      movieId      rating      timestamp
## Min.   :    1  Min.   :    1  Min.   :0.500  Min.   :7.897e+08
## 1st Qu.:18122  1st Qu.:   648  1st Qu.:3.000  1st Qu.:9.468e+08
## Median :35743  Median :  1834  Median :4.000  Median :1.035e+09
## Mean   :35869  Mean   :  4120  Mean   :3.512  Mean   :1.033e+09
## 3rd Qu.:53602  3rd Qu.:  3624  3rd Qu.:4.000  3rd Qu.:1.127e+09
## Max.   :71567  Max.   :65133  Max.   :5.000  Max.   :1.231e+09
##      title      genres
## Length:9000047  Length:9000047
## Class :character Class :character
## Mode  :character Mode  :character
##
##
##
```

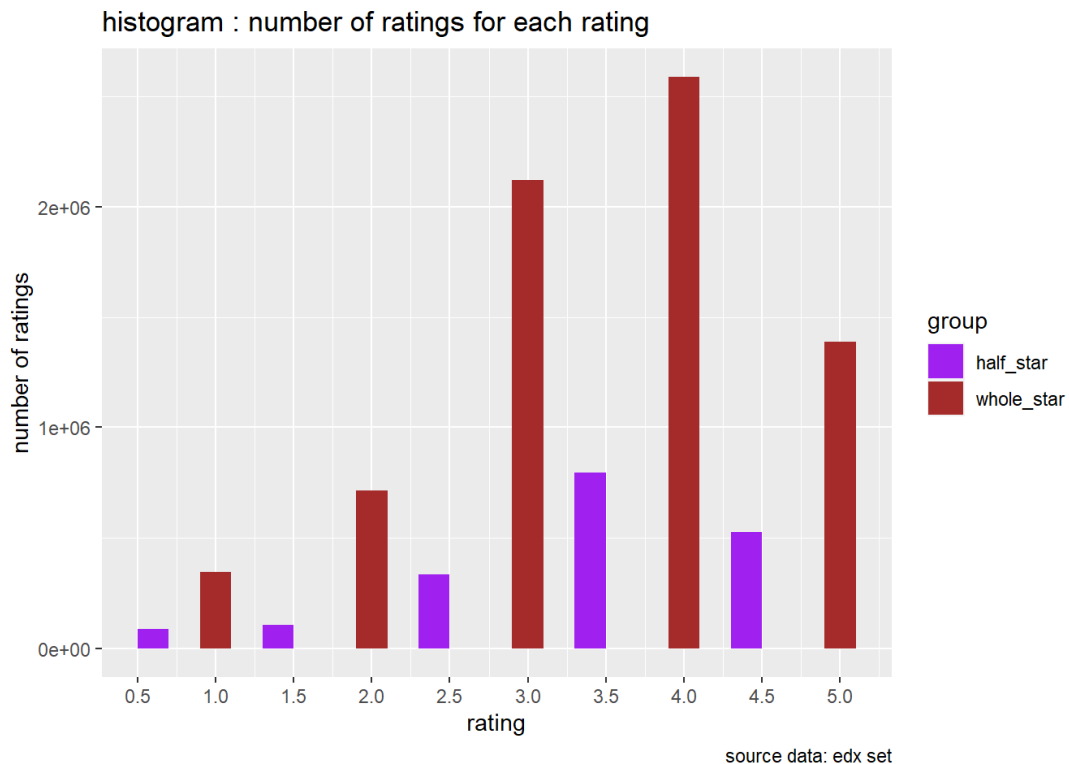
```
#A new dataframe "explore_ratings" is created which contains half star and whole star ratings from the edx set :
```

```
group <- ifelse((edx$rating == 1 | edx$rating == 2 | edx$rating == 3 |
  edx$rating == 4 | edx$rating == 5) ,
  "whole_star",
  "half_star")
```

```
explore_ratings <- data.frame(edx$rating, group)
```

```
# Plot the explore_ratings dataframe via histogram
```

```
ggplot(explore_ratings, aes(x= edx.rating, fill = group)) +  
  geom_histogram( binwidth = 0.2) +  
  scale_x_continuous(breaks=seq(0, 5, by= 0.5)) +  
  scale_fill_manual(values = c("half_star"="purple", "whole_star"="brown")) +  
  labs(x="rating", y="number of ratings", caption = "source data: edx set") +  
  ggtitle("histogram : number of ratings for each rating")
```



Exploring ratings of the edx set , we notice the following facts:

1. The average user's activity reveals that no user gives 0 as rating
2. The top 5 ratings from most to least are : 4, 3, 5, 3.5 and 2.
3. The histogram shows that the half star ratings are less common than whole star ratings.


Exploring the features "genres" and "title" of our edx set.

```
#bar chart of top_title
```

```
top_title <- edx %>%  
  group_by(title) %>%  
  summarize(count=n()) %>%  
  top_n(20,count) %>%  
  arrange(desc(count))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)

top_title %>%
  ggplot(aes(x=reorder(title, count), y=count)) +
  geom_bar(stat='identity', fill="blue") + coord_flip(y=c(0, 40000)) +
  labs(x="", y="Number of ratings") +
  geom_text(aes(label= count), hjust=-0.1, size=3) +
  labs(title="Top 20 movies title based \n on number of ratings" , caption =
"source data: edx set")
```



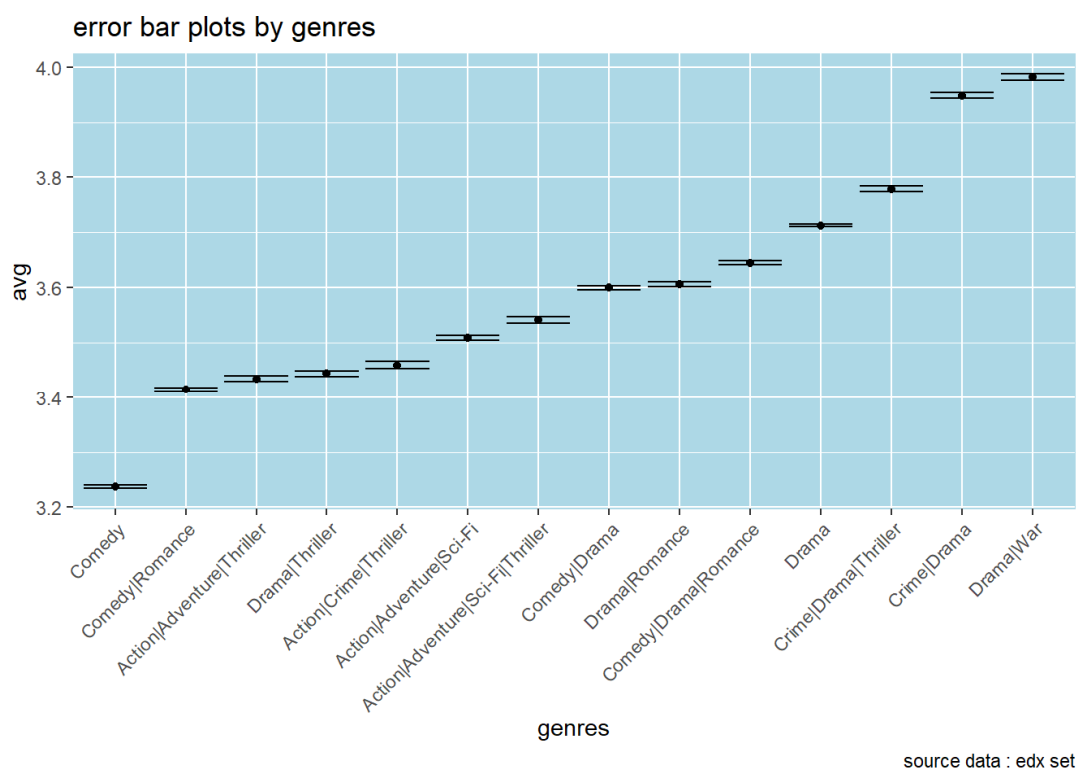
The movies which have the highest number of ratings are in the top genres categories : movies like Pulp fiction (1994), Forrest Gump(1994) or Jurassic Park(1993) which are in the top 5 of movie's ratings number , are part of the Drama, Comedy or Action genres.

#Computing the average and standard error for each "genre" , plotting the effect of genre

```
edx %>% group_by(genres) %>%
  summarize(n = n(), avg = mean(rating), se = sd(rating)/sqrt(n())) %>%
  filter(n >= 100000) %>%
  mutate(genres = reorder(genres, avg)) %>%
  ggplot(aes(x = genres, y = avg, ymin = avg - 2*se, ymax = avg + 2*se)) +
  geom_point() +
  geom_errorbar() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs(title = "error bar plots by genres" , caption = "source data : edx set")
```

```
) +
theme(
  panel.background = element_rect(fill = "lightblue",
                                   colour = "lightblue",
                                   size = 0.5, linetype = "solid"),
  panel.grid.major = element_line(size = 0.5, linetype = 'solid',
                                   colour = "white"),
  panel.grid.minor = element_line(size = 0.25, linetype = 'solid',
                                   colour = "white")
)

## `summarise()` ungrouping output (override with `.groups` argument)
```



We observe that the generated plot shows strong evidence of a genre effect.


```
edx %>%
  summarize(n_users = n_distinct(userId),
            n_movies = n_distinct(movieId))
```

Even if each row represents a rating given by one user to one movie, the number of unique values for the userId is 69878 and for the movieId 10664 : Both userId and movieId which are presented as integer should be presumably treat as factors for some analysis purposes. Also, this means that there are less movies provided for ratings than users that rated them. If we think in terms of a large matrix, with user on the rows and movies on the columns, a challenge we face is the sparsity of our matrix. This

large matrix will contain many empty cells. Moreover, we face a curse of dimensionality problem. These issues should be treated in our further analysis.

```
# histogram of number of ratings by movieId
```


```
edx %>%
  count(movieId) %>%
  ggplot(aes(n)) +
  geom_histogram( bins=30, color = "red") +
  scale_x_log10() +
  ggtitle("Movies") +
  labs(subtitle = "number of ratings by movieId",
       x="movieId" ,
       y="number of ratings",
       caption = "source data : edx set") +
  theme(panel.border = element_rect(colour="black", fill=NA))
```



source data : edx set

```
edx %>%
  count(userId) %>%
  ggplot(aes(n)) +
  geom_histogram( bins=30, color = "red") +
  scale_x_log10() +
  ggtitle("Users") +
  labs(subtitle = "number of ratings by UserId",
       x="userId" ,
```

```
y="number of ratings") +
theme(panel.border = element_rect(colour="black", fill=NA))
```



DATA PREPROCESSING Data typically needs to be preprocessed (e.g. cleansed, filtered, transformed) in order to be used by the machine learning techniques in the analysis step.

1. Data transformation: Building a rating matrix

#Using SparseMatrix function to get the rating matrix from Matrix package
library(Matrix)

```
##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack

edx_1 <- edx

edx_1$userId <- as.factor(edx_1$userId)
edx_1$movieId <- as.factor(edx_1$movieId)

edx_1$userId <- as.numeric(edx_1$userId)
edx_1$movieId <- as.numeric(edx_1$movieId)
sparse_ratings <- sparseMatrix(i = edx_1$userId,
                               j = edx_1$movieId ,
```

```

x = edx_1$rating,
dims = c(length(unique(edx_1$userId)),
          length(unique(edx_1$movieId))),
dimnames = list(paste("u", 1:length(unique(edx_1$use
rId))), sep = ""),
                paste("m", 1:length(unique(edx_1$movi
eId))), sep = "")))

```

```

# remove the copy created
rm(edx_1)

```

```

#give a look on the first 10 users
sparse_ratings[1:10,1:10]

```

```

## 10 x 10 sparse Matrix of class "dgCMatrix"

```

```

## [[ suppressing 10 column names 'm1', 'm2', 'm3' ... ]]

```

```

##
## u1 . . . . .
## u2 . . . . .
## u3 . . . . .
## u4 . . . . .
## u5 1 . . . . 3 . . .
## u6 . . . . .
## u7 . . . . .
## u8 . 2.5 . . 3 4 . . .
## u9 . . . . .
## u10 . . . . . 3 . . .

```

```

class(sparse_ratings)

```

```

## [1] "dgCMatrix"
## attr(,"package")
## [1] "Matrix"

```

```

#Convert rating matrix into a recommenderlab sparse matrix via recommenderlab
package

```

```

library(recommenderlab)

```

```

ratingMat <- new("realRatingMatrix", data = sparse_ratings)
ratingMat

```

```

## 69878 x 10664 rating matrix of class 'realRatingMatrix' with 9000047 ratings.

```

2. Relevant Data

We know that some users saw more movies than the others. So, instead of displaying some random users and movies, we should select the most relevant users and movies. Thus we visualize only the users who have seen many movies and the movies that have been seen by many users. To identify and select the most relevant users and movies, we follow these steps:

1. Determine the minimum number of movies per user.
2. Determine the minimum number of users per movie.
3. Select the users and movies matching these criteria.

```
min_n_movies <- quantile(rowCounts(ratingMat), 0.9)
```

```
min_n_users <- quantile(colCounts(ratingMat), 0.9)
```

```
ratings_movies <- ratingMat[rowCounts(ratingMat) > min_n_movies,  
                             colCounts(ratingMat) > min_n_users]
```

we can notice that now, we have a rating matrix of 6976 distinct users (rows) x 1067 distinct movies(columns) , with 2311476 ratings .

```
#before to proceed with regularization, i just remove the object copy of validation, "valid"
```

```
rm(valid)
```

```
## Warning in rm(valid): object 'valid' not found
```

```
#e. regularization
```

```
# remembering (5),  $\lambda$  is a tuning parameter. We can use cross-validation to choose it
```

```
lambdas <- seq(0, 10, 0.25)
```

```
rmses <- sapply(lambdas, function(l){
```

```
  mu_reg <- mean(edx$rating)
```

```
  b_i_reg <- edx %>%  
    group_by(movieId) %>%  
    summarize(b_i_reg = sum(rating - mu_reg)/(n()+1))
```

```
  b_u_reg <- edx %>%  
    left_join(b_i_reg, by="movieId") %>%  
    group_by(userId) %>%  
    summarize(b_u_reg = sum(rating - b_i_reg - mu_reg)/(n()+1))
```

```
  predicted_ratings_b_i_u <-  
    validation %>%
```


```

left_join(b_i_reg, by = "movieId") %>%
left_join(b_u_reg, by = "userId") %>%
mutate(pred = mu_reg + b_i_reg + b_u_reg) %>%
.$pred

return(RMSE(validation$rating,predicted_ratings_b_i_u))
})

qplot(lambdas, rmses)

```



```

lambda <- lambdas[which.min(rmses)]
lambda

## [1] 5.5

#valid_set
mu <- mean(edx$rating)
b_i_reg <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+lambda))

## `summarise()` ungrouping output (override with `.groups` argument)

b_u_reg <- edx %>%
  left_join(b_i_reg, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lambda))

```

```
## `summarise()` ungrouping output (override with `.groups` argument)

predicted_ratings_6 <-
  validation %>%
    left_join(b_i_reg, by = "movieId") %>%
    left_join(b_u_reg, by = "userId") %>%
    mutate(pred = mu + b_i + b_u) %>%
    pull(pred)
View(predicted_ratings_6)
model_6_rmse <- RMSE(predicted_ratings_6, validation$rating) # 0.864818
```

Methods and Analysis Recommender Engines

a. POPULAR , UBCF and IBCF algorithms of the recommenderlab package

```
library(recommenderlab)
```

```
model_pop <- Recommender(ratings_movies, method = "POPULAR",
  param=list(normalize = "center"))
```

#prediction example on the first 10 users

```
pred_pop <- predict(model_pop, ratings_movies[1:10], type="ratings")
as(pred_pop, "matrix")[,1:10]
```

```
##           m1           m2           m3           m5           m6           m7           m9           m
10
## u8      3.845709           NA 2.908521           NA           NA 3.091285 2.524774 3.3150
59
## u17           NA           NA 3.012617           NA 3.860392           NA 2.628869
NA
## u28           NA           NA           NA           NA 3.167586 2.502576           NA 2.7263
50
## u30           NA           NA           NA 2.800736           NA 3.118312 2.551801
NA
## u43      4.693015 3.793565 3.755826 3.621014 4.603602           NA 3.372079 4.1623
65
## u48           NA           NA           NA 3.505851 4.488439 3.823428 3.256917
NA
## u57           NA           NA 2.646955 2.512143 3.494730 2.829720 2.263208 3.0534
93
## u70      4.426355 3.526905 3.489166 3.354354 4.336941 3.671931 3.105419 3.8957
04
## u88           NA 3.040854 3.003116 2.868303           NA 3.185880 2.619368 3.4096
54
## u103           NA 2.819650 2.781912 2.647099           NA 2.964676 2.398164 3.1884
50
##           m11           m14
## u8      3.452896 3.411770
## u17           NA 3.515866
## u28      2.864186 2.823060
## u30           NA 3.438797
## u43           NA 4.259075
## u48      4.185039 4.143913
```

```
## u57          NA 3.150204
## u70  4.033541 3.992415
## u88  3.547490 3.506365
## u103 3.326286      NA
```

```
#Calculation of rmse for popular method
```

```
e <- evaluationScheme(ratings_movies, method="split", train=0.7, given=-5)
#5 ratings of 30% of users are excluded for testing
```

```
model_pop <- Recommender(getData(e, "train"), "POPULAR")
```

```
prediction_pop <- predict(model_pop, getData(e, "known"), type="ratings")
```

```
rmse_popular <- calcPredictionAccuracy(prediction_pop, getData(e, "unknown"))
[1]
rmse_popular
```

```
##          RMSE
## 0.8470042
```

```
#Estimating rmse for UBCF using Cosine similarity and selected n as 50 based
on cross-validation
```

```
set.seed(1)
```

```
model <- Recommender(getData(e, "train"), method = "UBCF",
                     param=list(normalize = "center", method="Cosine", nn=50)
)
```

```
prediction <- predict(model, getData(e, "known"), type="ratings")
```

```
rmse_ubcf <- calcPredictionAccuracy(prediction, getData(e, "unknown"))[1]
rmse_ubcf
```

```
##          RMSE
## 0.8320826
```

```
#Estimating rmse for IBCF using Cosine similarity and selected n as 350 based
on cross-validation
```

```
set.seed(1)
```

```
model_ibcf <- Recommender(getData(e, "train"), method = "IBCF",
                          param=list(normalize = "center", method="Cosine", k=350)
)
```

```
prediction_ibcf <- predict(model_ibcf, getData(e, "known"), type="ratings")
```

```
rmse_ibcf <- calcPredictionAccuracy(prediction_ibcf, getData(e, "unknown"))[1]
rmse_ibcf
```

```
##      RMSE
## 0.9569868
```

```
#summarize all the rmse for recommender algorithms
library(kableExtra)
```

```
rmse_results <- data.frame(methods=c("Regularized Movie + User Effect Model",
"Recommender Popular Model" , "Recommender UBCF" ,"Recommender IBCF"),rmse =
c(model_6_rmse, rmse_popular,rmse_ubcf,rmse_ibcf))
```

```
kable(rmse_results) %>%
  kable_styling(bootstrap_options = "striped" , full_width = F , position = "
center") %>%
  kable_styling(bootstrap_options = "bordered", full_width = F , position ="c
enter") %>%
  column_spec(1,bold = T ) %>%
  column_spec(2,bold = T ,color = "white" , background ="#D7261E")
```

Methods	RMSE
Regularized Movie + User Effect Model	0.8649855
Recommender Popular Model	0.8470042
Recommender UBCF	0.8320826
Recommender IBCF	0.9569868